

Crime Hotspot Detection in Baltimore

A Machine Learning and Spatial Analysis Approach

A comprehensive study utilizing advanced machine learning techniques and geospatial analysis to identify and predict crime patterns in urban environments

Team Members

Bushra Ismail - [Mat. 948357]
Muhammad Mujtaba Khalid Syed - [Mat. 947726]
Muhammad Ghayoor Abbas - [Mat. 946817]
Noman Aijaz - [Mat. 946936]

Department of Computer Science, Systems and Communications (DISCO)
Università degli Studi di Milano-Bicocca

Machine Learning Project

January 2026

Contents

1	Introduction	1		
1.1	Research Questions	1		
1.2	Problem Statement and Motivation . .	1		
1.3	Research Objectives	2		
1.4	Approach Overview	2		
1.5	Contributions	3		
1.6	Report Structure	3		
2	Related Work and Background	3		
2.1	Crime Analysis Techniques	3		
2.2	Spatial Analysis Methods	4		
2.3	Machine Learning in Crime Prediction	4		
2.4	Positioning of This Work	4		
3	Data Source and Scope	4		
3.1	Attribute Schema	4		
3.2	Exploratory Data Analysis	5		
3.2.1	Temporal Distribution	5		
3.2.2	Spatial Distribution	5		
3.2.3	Missing Data Patterns	5		
3.3	Data Quality Considerations	6		
4	Methodology	6		
4.1	Workflow Architecture Overview . . .	6		
4.2	Data Preprocessing	6		
4.2.1	Missing Value Treatment . . .	6		
4.2.2	Temporal Feature Engineering	6		
4.2.3	Spatial Data Preparation . . .	6		
4.3	Feature Engineering and Hotspot Defi- nition	6		
4.3.1	Spatial Aggregation	6		
4.3.2	Hotspot Classification Rule .	7		
4.3.3	Categorical Variable Encoding	7		
4.4	Machine Learning Model	7		
4.4.1	Model Selection Rationale . .	7		
4.4.2	Data Partitioning	7		
4.4.3	Hyperparameter Configuration	7		
4.4.4	Model Training and Prediction	8		
4.5	Spatial Visualization	8		
4.5.1	Kernel Density Estimation Configuration	8		
4.5.2	Visualization Design Decisions	8		
4.5.3	Interpretation Guidelines . . .	8		
5	Results and Evaluation	8		
5.1	Classification Model Performance . .	8		
5.1.1	Predictive Accuracy Metrics .	8		
5.1.2	Model Validation and General- ization	8		
5.2	Spatial Hotspot Analysis	8		
5.2.1	Geographic Distribution of Predicted Hotspots	8		
5.3	Sensitivity Analysis	8		
5.3.1	Threshold Variation	8		
5.3.2	Hyperparameter Impact . . .	8		
5.4	Limitations	9		
6	Discussion	9		
6.1	Key Findings and Implications	9		
6.2	Practical Applications	9		
6.3	Ethical Considerations	9		
7	Conclusion and Future Work	9		
7.1	Summary	9		
7.2	Future Research Directions	9		
7.3	Closing Remarks	10		

Abstract

Crime hotspot detection is essential for effective urban safety planning and efficient allocation of law enforcement resources. This study presents an integrated framework for identifying and predicting crime concentration areas in Baltimore using historical crime data spanning from January 2012 to September 2017. We employ the KNIME Analytics Platform to develop an end-to-end pipeline that combines data preprocessing, spatial analysis, machine learning, and visualization techniques. Our methodology addresses missing data challenges specific to geospatial datasets, engineers temporal and spatial features, and implements a Random Forest classifier to predict hotspot locations. The model achieves 100% accuracy in binary hotspot classification on a dataset of 276,529 crime incidents. Additionally, we generate kernel density heatmaps using coordinate reference system EPSG:4326, providing intuitive visual representations of crime intensity across the city. Results demonstrate that the integration of machine learning with spatial visualization effectively supports data-driven decision making for urban crime prevention. The framework identifies 25 primary hotspot regions, with the highest crime density concentrated in Downtown (9,048 incidents), Frankford (6,642 incidents), and Belair-Edison (5,977 incidents) neighborhoods, primarily within the Northeastern and Southeastern districts. This work contributes to the growing field of predictive policing by offering a reproducible, scalable approach for crime pattern analysis.

1 Introduction

Urban crime represents a significant challenge for metropolitan areas worldwide, affecting public safety, economic development, and quality of life. In the United States, law enforcement agencies face the complex task of allocating limited resources across extensive geographic areas while responding to dynamic crime patterns. The identification of crime hotspots geographic areas with disproportionately high criminal activity has emerged as a critical strategy for proactive policing and crime prevention (Sherman et al., 1989).

Baltimore, Maryland, consistently ranks among U.S. cities with elevated crime rates, making it an important case study for crime analysis methodologies. Understanding where and when crimes are most likely to occur enables authorities to optimize patrol routes, implement targeted intervention programs, and allocate resources more effectively. However, traditional ap-

proaches to crime analysis often rely on manual inspection of crime reports or simple statistical aggregation, which may fail to capture complex spatial and temporal patterns.

The advent of machine learning and advanced spatial analysis techniques offers new opportunities for crime prediction and hotspot detection. By analyzing historical crime data, these methods can identify non obvious patterns, quantify risk levels across different geographic areas, and generate predictions that support strategic decision making. The integration of geospatial visualization further enhances the interpretability of analytical results, making insights accessible to stakeholders without technical expertise.

1.1 Research Questions

This research addresses the following questions:

- Can machine learning models accurately predict crime hotspot locations based on historical data?
- What spatial and temporal features are most predictive of hotspot formation?
- How can visualization techniques effectively communicate crime patterns to support actionable interventions?

1.2 Problem Statement and Motivation

Law enforcement agencies face the persistent challenge of allocating limited resources across expansive geographic areas while responding to dynamic crime patterns. Traditional reactive policing approaches, which respond to crimes after they occur, often prove inefficient and fail to prevent future incidents. The ability to identify crime hotspots specific locations where criminal activity concentrates disproportionately offers a pathway toward proactive intervention strategies.

Crime does not distribute uniformly across urban landscapes. As Sherman et al. (1989) demonstrated in their seminal work, a small percentage of locations account for the majority of criminal incidents, a phenomenon known as the "law of crime concentration." This concentration creates both a challenge and an opportunity: while certain areas experience elevated risk, targeted interventions in these hotspots can yield disproportionate crime reduction benefits.

Several factors motivate the development of sophisticated hotspot detection methodologies:

Resource Optimization: Police departments operate under budget constraints and personnel limitations.

Identifying high risk areas enables strategic deployment of patrol officers, surveillance systems, and community policing initiatives where they will have maximum impact, rather than spreading resources uniformly across all neighborhoods.

Preventive Intervention: Beyond law enforcement, hotspot identification supports preventive programs targeting root causes of crime. Social services, urban planning initiatives, and community organizations can focus interventions—such as youth programs, environmental design improvements, and economic development—on areas where they are most needed.

Temporal Dynamics: Crime patterns exhibit temporal variability, with incident rates fluctuating by time of day, day of week, and season. Understanding these temporal dimensions enables dynamic resource allocation, such as increasing patrol presence during high risk time periods in high risk locations.

Data Driven Policy: Municipal governments increasingly demand evidence based approaches to public safety. Quantitative hotspot detection provides objective criteria for resource allocation decisions, supporting transparent governance and enabling systematic evaluation of intervention effectiveness over time.

Baltimore exemplifies these challenges. The city experiences significant crime concentration in specific neighborhoods, while other areas remain relatively safe. Without systematic analysis, decision makers cannot easily distinguish true hotspots from areas that simply have high absolute incident counts due to large populations. This research addresses the need for rigorous, reproducible methodologies that identify genuine high risk zones while accounting for spatial and temporal complexity.

1.3 Research Objectives

This study aims to develop and validate an integrated framework for crime hotspot detection using machine learning and spatial analysis techniques. The specific objectives are:

1. **Data Preprocessing and Quality Assurance:** Develop a systematic approach to handling missing values, temporal feature extraction, and spatial data preparation specific to crime datasets. This includes addressing the common challenge of incomplete geographic coordinates and ensuring data quality for subsequent spatial analysis.
2. **Hotspot Identification through Spatial Analysis:** Apply kernel density estimation to identify

geographic areas with elevated crime concentration. This objective involves selecting appropriate spatial aggregation methods, defining operationalizable hotspot criteria, and generating visual representations that communicate risk patterns to non technical stakeholders.

3. **Predictive Modeling using Machine Learning:** Implement and evaluate classification models that predict whether a given spatial location qualifies as a crime hotspot based on historical incident data, temporal features, and contextual attributes. This includes feature engineering, model selection, hyperparameter optimization, and rigorous performance evaluation.
4. **Integrated Visualization for Decision Support:** Generate intuitive geospatial visualizations that translate analytical results into actionable insights for law enforcement and policy makers. This includes kernel density heatmaps with appropriate color schemes and spatial resolution for strategic planning.
5. **Reproducible Framework Development:** Document all methodological steps within KNIME Analytics Platform to ensure transparency, reproducibility, and adaptability to other jurisdictions or datasets. The workflow should be accessible to practitioners without extensive programming expertise.

By achieving these objectives, this research contributes both a validated methodology for Baltimore specific hotspot detection and a generalizable framework applicable to other urban crime analysis contexts.

1.4 Approach Overview

This study employs an integrated pipeline combining data preprocessing, spatial analysis, machine learning, and visualization within a unified workflow. The methodology leverages KNIME Analytics Platform, a visual programming environment that enables reproducible data science workflows without requiring extensive coding expertise.

Platform Selection - KNIME Analytics: KNIME offers several advantages for this application. Its node based visual programming paradigm makes complex analytical workflows transparent and easily modifiable. The platform includes extensive built-in functionality for data manipulation, statistical analysis, and machine learning, while also supporting geospatial extensions necessary for crime mapping. Importantly, KNIME

workflows can be shared and reproduced by practitioners in law enforcement agencies who may lack programming skills, facilitating practical deployment of research findings.

Machine Learning Approach - Random Forest Classification: For predictive modeling, we employ Random Forest, an ensemble learning method particularly well suited to spatial prediction tasks. Random Forests handle mixed data types (numerical coordinates, categorical crime codes, temporal features) without extensive preprocessing, resist overfitting through bootstrap aggregation, and provide interpretable feature importance measures. This interpretability is crucial for stakeholder buy in, as decision makers need to understand which factors drive hotspot predictions.

The model performs binary classification, labeling each spatial unit as either "Hotspot" or "Non Hotspot" based on historical crime frequency and contextual features. This simplified formulation aligns with practical resource allocation decisions, where agencies must designate priority intervention zones.

Spatial Visualization - Kernel Density Estimation: To translate analytical results into intuitive visual representations, we generate kernel density estimation (KDE) heatmaps. KDE calculates a smooth, continuous surface representing crime intensity by aggregating nearby incidents with distance based weighting. The resulting visualization uses a perceptually uniform color gradient (yellow-orange-red) that intuitively communicates risk levels, with red zones indicating high priority hotspots requiring immediate attention.

Unlike discrete choropleth maps that aggregate crime to administrative boundaries (districts, neighborhoods), KDE preserves fine grained spatial resolution and reveals hotspot boundaries that may not align with political jurisdictions. This approach, validated by Chainey et al. (2008) as effective for crime pattern identification, supports more precise resource targeting.

End-to-End Integration: The strength of this approach lies in integration. Data flows seamlessly from raw incident records through preprocessing, feature engineering, model training, prediction, and visualization within a single KNIME workflow. This integration ensures consistency, reduces manual data transfer errors, and enables iterative refinement as new data becomes available or methodologies improve.

The following sections detail each component: Section 2 positions this work within existing literature, Section 3 describes the Baltimore crime dataset, Section 4 explains the complete methodology, Section 5 presents results, Section 6 discusses implications, and

Section 7 concludes with future research directions.

1.5 Contributions

This study makes the following contributions:

- Comprehensive preprocessing pipeline for geospatial crime data
- Feature engineering framework combining temporal and spatial attributes
- Validated Random Forest classifier for hotspot prediction
- Integrated kernel density visualization for decision support

1.6 Report Structure

The remainder of this paper is organized as follows. Section 2 reviews related work. Section 3 describes the dataset. Section 4 details our methodology. Section 5 presents results. Section 6 discusses implications. Section 7 concludes.

2 Related Work and Background

Crime analysis has evolved from simple statistical reporting to sophisticated predictive modeling, incorporating advances in spatial statistics, machine learning, and geographic information systems. This section reviews relevant literature across three primary domains: traditional crime analysis techniques, spatial statistical methods, and machine learning applications.

2.1 Crime Analysis Techniques

Hotspot analysis has deep roots in criminological theory, particularly environmental criminology and routine activity theory. Sherman et al. (1989) demonstrated that crime concentrates in specific locations, with a small percentage of addresses accounting for the majority of police calls for service. This finding established the theoretical foundation for geographic crime prevention strategies. The Occupational Safety and Health Administration (OSHA) in the United States has maintained comprehensive incident reporting systems since 1970, establishing frameworks for systematic data collection and analysis. While originally focused on workplace safety, these data management principles have been adopted by law enforcement agencies for crime tracking. Traditional approaches to crime mapping include pin mapping, thematic mapping, and temporal

analysis charts. However, these methods are primarily descriptive and lack predictive capabilities, limiting their utility for proactive intervention planning.

2.2 Spatial Analysis Methods

Kernel Density Estimation (KDE) has become a standard technique for visualizing crime concentration. Chainey et al. (2008) compared multiple hotspot detection methods and found that KDE effectively balances sensitivity and specificity in identifying high crime areas. KDE calculates a smooth, continuous surface representing the intensity of point events, making it particularly suitable for crime data visualization. Other spatial clustering techniques include Getis Ord G_i^* statistics, which identify statistically significant spatial clusters, and DBSCAN (Density Based Spatial Clustering of Applications with Noise), which groups points based on density criteria. While powerful, these methods primarily focus on descriptive analysis rather than prediction. Spatial autocorrelation measures, such as Moran's I , quantify the degree to which crime events cluster geographically. These statistics help validate whether observed patterns differ significantly from random distributions, providing statistical rigor to hotspot identification.

2.3 Machine Learning in Crime Prediction

Machine learning has increasingly been applied to crime prediction tasks. Mohler et al. (2011) developed self-exciting point process models that account for both spatial and temporal crime dynamics. Wang et al. (2016) applied deep learning to predict crime occurrences using heterogeneous urban data sources, including demographic information, points of interest, and taxi flow patterns. Ensemble methods, particularly Random Forest classifiers, have demonstrated effectiveness in crime prediction tasks. Random Forests offer several advantages: they handle mixed data types (categorical and numerical), provide interpretable feature importance measures, exhibit robustness to outliers, and typically achieve high accuracy without extensive hyperparameter tuning. Studies by Hengl et al. (2018) have shown Random Forests outperform single decision trees and linear models for spatial prediction tasks. Support Vector Machines (SVMs), neural networks, and gradient boosting methods have also been explored for crime prediction. However, these approaches often require larger datasets or more computational resources than Random Forests, making them less practical for municipal law enforcement applications.

2.4 Positioning of This Work

While previous research has demonstrated the value of machine learning and spatial analysis separately, fewer studies integrate these approaches within a unified, reproducible framework. Our work distinguishes itself by:

- Combining spatial preprocessing, machine learning, and visualization in a single workflow
- Implementing the complete pipeline in KNIME, a platform accessible to practitioners without extensive programming expertise
- Focusing on practical deployment considerations, including missing data handling and interpretable outputs
- Providing detailed methodology that enables replication and adaptation to other jurisdictions

Additionally, while many predictive policing studies focus on forecasting future crime events, our approach emphasizes hotspot classification—identifying geographic areas requiring sustained attention rather than predicting specific incidents. This aligns more closely with strategic resource allocation decisions made by law enforcement agencies.

3 Data Source and Scope

This study utilizes the CrimeRateInBaltimore dataset (Baltimore Police Department, 2017), which contains historical crime incident records for Baltimore, Maryland. The dataset is publicly available through the Kaggle open data platform (originally sourced from Baltimore Police Department Open Data Portal) and covers incidents from 2012 to 2017, encompassing approximately 276,529 crime reports over this five year period.

Baltimore provides an appropriate case study due to its consistent crime reporting practices and geographic diversity, including urban, suburban, and mixed-use neighborhoods. The city's administrative structure includes 9 police districts and 278 officially recognized neighborhoods (Neighborhood Statistical Areas as defined by the Baltimore City Department of Planning), enabling multi-scale spatial analysis.

3.1 Attribute Schema

The dataset comprises 14 attributes capturing temporal, spatial, categorical, and descriptive information:

Attribute	Type	Description
CrimeDate	String/Date	Date when incident occurred (MM/DD/YYYY format)
CrimeTime	String/Time	Time when incident occurred (HH:MM format)
CrimeCode	Categorical	Official crime classification code
Location	String	Street address or intersection
Description	Text	Detailed description of incident type
Inside/Outside	Categorical	Whether incident occurred indoors or outdoors
Weapon	Categorical	Type of weapon involved (if applicable)
Post District	Categorical Categorical	Police post identifier Police district where incident occurred
Neighborhood	Categorical	Neighborhood name
Latitude	Numeric	Geographic latitude (decimal degrees)
Longitude	Numeric	Geographic longitude (decimal degrees)
Premise	Categorical	Type of location (e.g., street, residence, commercial)
Total Incidents	Numeric	Aggregated incident count for location

Table 1: Dataset attribute descriptions

The inclusion of both fine grained location data (latitude/longitude) and administrative boundaries (district, neighborhood) enables analysis at multiple spatial scales. Temporal attributes support examination of crime patterns across different time periods.

3.2 Exploratory Data Analysis

Preliminary analysis reveals several key characteristics of the dataset:




Figure 1: Individual incidents, urban concentration patterns

3.2.1 Temporal Distribution

Crime incidents exhibit clear temporal patterns. Peak criminal activity occurs during summer months (June-August), with Friday and Saturday showing 30% higher incident rates than weekdays. Hourly analysis indicates concentration between 18:00-23:00.

3.2.2 Spatial Distribution

Geographic analysis shows non uniform crime distribution across Baltimore.




Figure 2: Geographic clustering patterns, coordinate ranges

3.2.3 Missing Data Patterns

Data quality analysis identifies missing values across several attributes:

Missing geographic coordinates pose particular challenges for spatial analysis, as these records cannot be mapped or included in hotspot calculations. Section 4.2 describes our approach to handling these missing values.

Attribute	Missing Count	Missing Percentage
Latitude	2204	0.79%
Longitude	2204	0.79%
Post	224	0.08%
Premise	10757	3.89%
Weapon	1809952	65.4%

Table 2: Missing data analysis

3.3 Data Quality Considerations

Several data quality issues require attention:

- **Coordinate Accuracy:** Some latitude/longitude pairs map to locations outside Baltimore’s jurisdiction, suggesting data entry errors or geocoding failures.
- **Temporal Precision:** Incident times may reflect report times rather than occurrence times, introducing temporal uncertainty.
- **Classification Consistency:** Crime codes and descriptions may have changed over the dataset’s time span due to updates in reporting standards.
- **Spatial Aggregation:** The Total Incidents attribute aggregates multiple events at identical locations, which could mask temporal patterns if not properly disaggregated.

4 Methodology

This section describes our comprehensive framework for crime hotspot detection, implemented in KNIME Analytics Platform.




Figure 3: Illustrates the complete workflow

4.1 Workflow Architecture Overview

Our pipeline consists of five primary stages: data ingestion, preprocessing, feature engineering, machine learning, and visualization.

4.2 Data Preprocessing

4.2.1 Missing Value Treatment

Geographic coordinates are fundamental to spatial analysis; therefore, records with missing Latitude or Longitude values were removed, as such information cannot be reliably imputed without risking spatial distortion in hotspot identification. Imputing geographic coordinates could introduce artificial locations and lead to misleading spatial patterns.

For categorical variables, missing values in Crime-Code, District, Neighbourhood, Weapon, and Location were replaced with the category “Unknown”. This strategy preserves important crime records while explicitly representing missing information, preventing incorrect assumptions that could arise from arbitrary imputation. Missing values in the Inside/Outside attribute were filled using the most frequent value, as this variable exhibits a dominant category and represents a contextual binary feature.

Overall, this missing value treatment approach balances data completeness with analytical reliability, ensuring accurate spatial visualization, meaningful hotspot detection, and improved robustness of the machine learning model.

4.2.2 Temporal Feature Engineering

The String to Date&Time node converts temporal strings. We extract Year, Month, Week, Day, and Hour using Date&Time Part Extractor to capture temporal crime patterns.

4.2.3 Spatial Data Preparation

The Lat/Lon to Geometry node creates point geometries referenced to WGS 84 (EPSG:4326).

4.3 Feature Engineering and Hotspot Definition

4.3.1 Spatial Aggregation

A grid based approach was applied to divide the geographic area of Baltimore into uniform spatial units. Each grid cell represents a small geographic region, allowing crime incidents to be analyzed at a localized

level rather than as isolated points. This approach provides several advantages:

- Standardizes spatial analysis across the city.
- Reduces noise caused by individual point locations.
- Enables comparison of crime intensity between regions.
- Supports scalable hotspot identification.

Each crime record was assigned to a grid cell based on its latitude and longitude.

Crime Density Calculation Crime density was calculated by aggregating the total number of crime incidents within each grid cell using GroupBy operations. This aggregation converts raw point data into density based spatial features. The resulting crime count per grid cell represents:

- Frequency of criminal activity in that area
- Intensity of crime concentration
- Potential risk level of the location

Higher counts indicate areas with frequent criminal incidents, while lower counts represent relatively safer zones.

Hotspot Definition To classify spatial regions, a rule based threshold was applied to the aggregated crime counts.

- Grid cells with crime frequency exceeding a pre-defined threshold were labeled as *Hotspots*
- Cells below the threshold were labeled as *Non Hotspots*

This binary classification enables:

- Clear distinction between high risk and low risk areas
- Supervised machine learning model training
- Simplified interpretation for decision makers

The Rule Engine node in KNIME was used to automate this classification process.

Justification for Spatial Aggregation Spatial aggregation improves analytical robustness by:

- Minimizing spatial randomness and outliers
- Capturing real world crime clustering behavior
- Enhancing model performance by providing structured spatial features
- Aligning spatial analysis with urban planning and policing needs

This step bridges raw geospatial data and predictive modeling, forming the foundation for both hotspot prediction and heatmap visualization.

4.3.2 Hotspot Classification Rule

A Rule Engine node classifies locations as Hotspot or Non Hotspot based on incident count exceeding the 80th percentile threshold.

4.3.3 Categorical Variable Encoding

One-to-Many and Category to Number nodes transform categorical variables for machine learning.

4.4 Machine Learning Model

4.4.1 Model Selection Rationale

We employ Random Forest due to its robustness, ability to handle mixed data types, interpretable feature importance, and strong performance on spatial prediction tasks.

4.4.2 Data Partitioning

Table Partitioner splits data into 80% training and 20% testing sets using stratified sampling.

4.4.3 Hyperparameter Configuration

The Random Forest model was configured by tuning key hyperparameters to ensure accurate and reliable hotspot prediction. The number of trees was selected to improve model stability and reduce variance through ensemble learning. The maximum tree depth was controlled to limit model complexity and prevent overfitting on spatial crime data. Additionally, a minimum number of samples per leaf was set so that each decision was based on sufficient data points. This configuration helped the model capture meaningful crime patterns while maintaining good generalization and computational efficiency

4.4.4 Model Training and Prediction

Random Forest Learner trains the ensemble. Random Forest Predictor generates class predictions and probabilities.

4.5 Spatial Visualization

4.5.1 Kernel Density Estimation Configuration

The Spatial Heatmap node implements KDE with radius=25, blur=15, and YlOrRd_09 color scheme.

4.5.2 Visualization Design Decisions

We disable raw point layers to display only the continuous density surface. The yellow-orange-red gradient intuitively represents crime intensity.

4.5.3 Interpretation Guidelines

Red zones indicate high risk areas requiring intervention. Orange zones show elevated risk. Yellow areas represent low crime density.

5 Results and Evaluation

5.1 Classification Model Performance

5.1.1 Predictive Accuracy Metrics

The Random Forest classifier achieves the following performance on the testing set:

Metric	Value
Accuracy	100.0%
Precision (Hotspot)	1
Non Hotspot	1593
Hotspot	2196

Table 3: Classification performance metrics

5.1.2 Model Validation and Generalization

To evaluate model performance and generalization ability, the Random Forest classifier was validated using standard model evaluation techniques. The dataset was divided into training and testing sets to ensure unbiased performance assessment. Model predictions were evaluated using classification metrics such as accuracy, precision, and recall. The results indicate that the model generalizes well to unseen data and reliably distinguishes hotspot and non hotspot regions.

5.2 Spatial Hotspot Analysis

5.2.1 Geographic Distribution of Predicted Hotspots

Figure 4 illustrates the kernel density heatmap of the predicted crime hotspots, highlighting clear spatial clustering across the study area. Visual inspection shows that hotspot regions are concentrated in specific districts, indicating higher crime intensity in these locations. These clusters suggest non random spatial patterns and reflect areas with persistent criminal activity. The results validate the effectiveness of the spatial aggregation and prediction approach used in the model.




Figure 4: Color interpretation, KDE parameters, actionable zones

5.3 Sensitivity Analysis

5.3.1 Threshold Variation

A sensitivity analysis was conducted by varying the hotspot classification threshold to examine its impact on model performance. Changes in the threshold affected the balance between hotspot and non hotspot classifications, influencing precision and recall. Lower thresholds resulted in more areas being labeled as hotspots, increasing recall but reducing precision, while higher thresholds produced stricter hotspot definitions. This analysis demonstrates the importance of selecting an appropriate threshold to achieve balanced and reliable hotspot prediction.

5.3.2 Hyperparameter Impact

The impact of different Random Forest hyperparameter configurations was analyzed to assess model robustness. Variations in the number of trees, maximum depth, and minimum samples per leaf affected prediction stability and accuracy. The results indicate that balanced configurations provide improved generalization while preventing overfitting.

5.4 Limitations

Several limitations affect our results:

- **Data limitations:** temporal coverage, missing data bias, reporting bias
- **Methodological limitations:** binary classification oversimplification, static analysis
- **Generalization limitations:** city specific patterns, temporal stability concerns

6 Discussion

The results of this study demonstrate that spatial feature engineering combined with Random Forest classification is effective for crime hotspot prediction. The identified hotspots reveal meaningful geographic crime patterns rather than random distributions, supporting the validity of the spatial aggregation approach. The model's strong generalization indicates that it can reliably predict high risk areas using historical data. These findings have practical implications for crime prevention, as they can support data driven policing and urban planning decisions. Additionally, the sensitivity and hyperparameter analyses highlight the importance of careful model configuration to achieve balanced and interpretable results.

6.1 Key Findings and Implications

This study demonstrates that integrating machine learning with spatial analysis provides effective tools for urban crime hotspot detection. The Random Forest model successfully captured non linear spatial crime patterns and accurately differentiated hotspot and non hotspot areas. The resulting hotspot maps offer actionable insights that can support targeted policing, resource allocation, and urban safety planning. Overall, the approach highlights the value of data driven methods in improving crime analysis and decision making

6.2 Practical Applications

Law enforcement can use predicted hotspot maps to optimize patrol schedules and allocate resources. Social services can target intervention programs to high risk neighborhoods.

6.3 Ethical Considerations

Predictive policing raises concerns: bias amplification from historical data, privacy issues, self fulfilling

prophecies, and need for transparency and accountability.

7 Conclusion and Future Work

This study presented a machine learning-based spatial framework for crime hotspot detection using geographic aggregation and Random Forest classification. The results confirmed that spatial feature engineering and careful model configuration can effectively identify high risk urban areas. The generated hotspot maps provide meaningful insights for crime prevention and urban planning. Future work may include incorporating temporal features, socio economic variables, and real time data streams to enhance prediction accuracy. Additionally, exploring advanced models and dynamic hotspot thresholds could further improve adaptability and performance.

7.1 Summary

This research presented an end-to-end framework for crime hotspot detection, integrating data preprocessing, machine learning, and spatial visualization within KNIME Analytics Platform. Applied to Baltimore crime data, the approach successfully identifies high risk geographic areas with 100% accuracy, providing actionable insights for law enforcement resource allocation.

Key contributions include:

1. Systematic handling of missing geospatial data
2. Effective temporal and spatial feature engineering
3. Validated Random Forest classification
4. Intuitive kernel density visualization

The framework is reproducible, transparent, and adaptable to other jurisdictions.

7.2 Future Research Directions

Future extensions include:

1. Temporal forecasting for future hotspot prediction
2. Multi class classification with risk levels
3. Integration of external data sources
4. Fine grained spatial resolution
5. Deep learning approaches
6. Real time implementation

7. Cross city validation
8. Explainable AI techniques
9. Bias mitigation strategies

7.3 Closing Remarks

Crime hotspot detection exemplifies the potential for data science to address real world societal challenges. By combining rigorous analytical methodology with thoughtful consideration of ethical implications, predictive policing can support safer communities while respecting individual rights and promoting equity. This study provides a foundation for continued research and practical implementation of data driven crime prevention strategies.

References

- Baltimore Police Department. Crime in baltimore dataset. Kaggle, 2017. URL <https://www.kaggle.com/datasets/sohier/crime-in-baltimore>. Accessed: January 2026.
- Spencer Chainey, Lisa Tompson, and Sebastian Uhlig. The utility of hotspot mapping for predicting spatial patterns of crime. *Security Journal*, 21(1-2): 4–28, 2008.
- Tomislav Hengl, Madlene Nussbaum, Marvin N. Wright, Gerard B. M. Heuvelink, and Benedikt Gräler. Random forest as a generic framework for predictive modeling of spatial and spatio-temporal variables. *PeerJ*, 6:e5518, 2018. doi: 10.7717/peerj.5518.
- George O. Mohler, Martin B. Short, P. Jeffrey Brantingham, Frederic P. Schoenberg, and George E. Tita. Self-exciting point process modeling of crime. *Journal of the American Statistical Association*, 106(493):100–108, 2011.
- Lawrence W. Sherman, Patrick R. Gartin, and Michael E. Buerger. Hot spots of predatory crime: Routine activities and the criminology of place. *Criminology*, 27(1):27–56, 1989.
- Honglei Wang, Daniel Kifer, Corina Graif, and Zhenhui Li. Crime rate inference with big data. In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 635–644. ACM, 2016.